COREY PIERSON

It's Not You, It's My Data

LEVERAGING CUSTOMER ANALYTICS TO BUILD A SCALABLE CUSTOMER CHURN PREVENTION SYSTEM

We miss you.

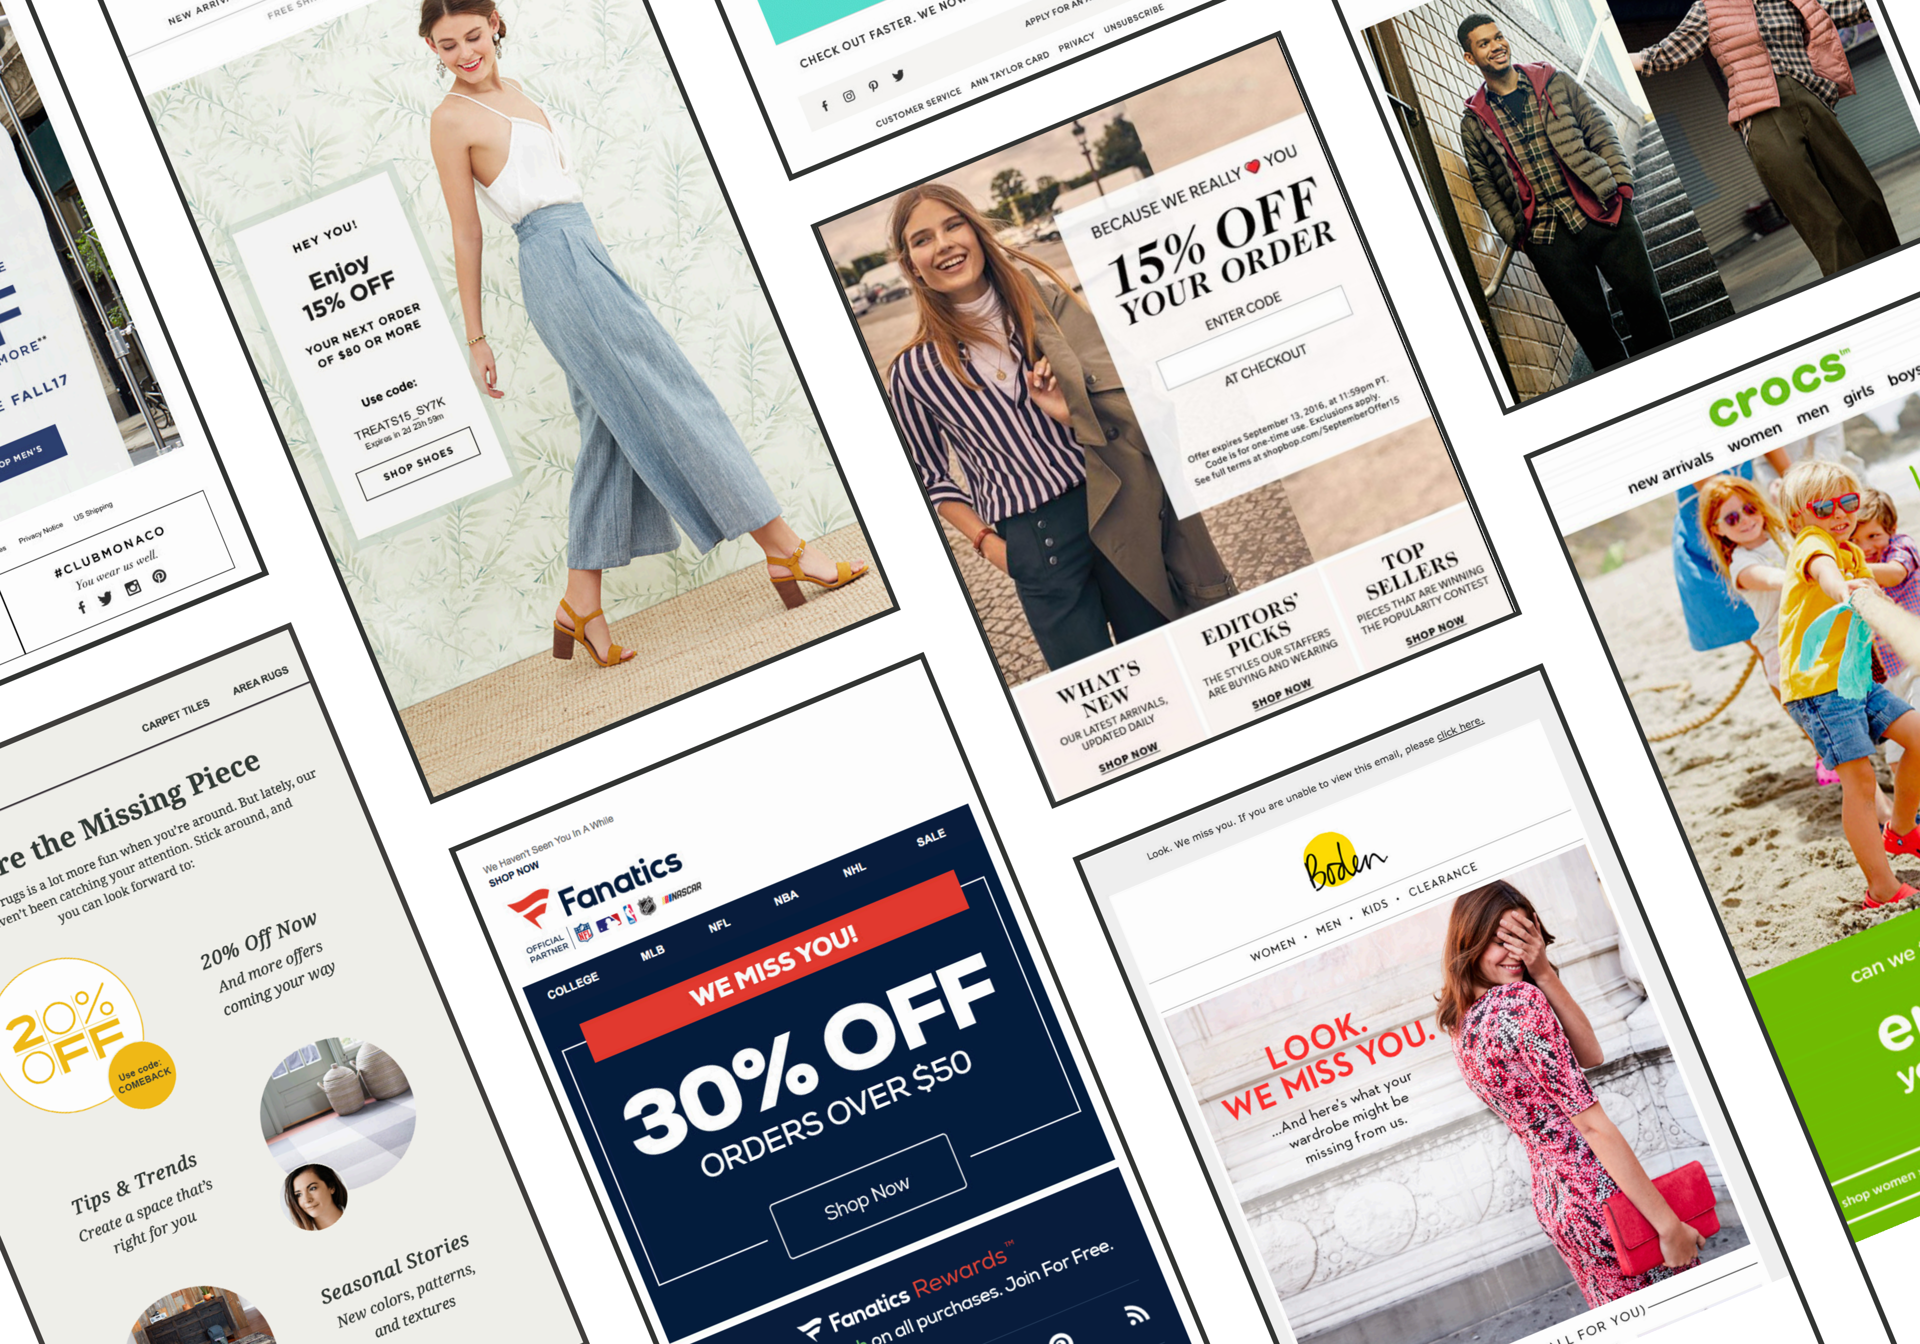
#CLUBMONACO
You wear us well.
HEY YOU!
Enjoy 15% OFF
YOUR NEXT ORDER OF $80 OR MORE
Use code:
TREATS15_SY7K
SHOP SHOES
CHECK OUT FASTER.
CUSTOMER SERVICE
ANN TAYLOR CARD
PRIVACY
UNSUBSCRIBE
BECAUSE WE REALLY ♥ YOU
15% OFF YOUR ORDER
ENTER CODE
AT CHECKOUT
Offer expires September 13, 2016, at 11:59pm PT. Code is for one-time use. Exclusions apply. See full terms at shopbop.com/SeptemberOffer15
WHAT'S NEW
OUR LATEST ARRIVALS, UPDATED DAILY
SHOP NOW
EDITORS' PICKS
THE STYLES OUR STAFFERS ARE BUYING AND WEARING
SHOP NOW
TOP SELLERS
PIECES THAT ARE WINNING THE POPULARITY CONTEST
SHOP NOW
crocs
new arrivals
women
men
girls
boys
CARPET TILES
AREA RUGS
the Missing Piece
20% Off Now
And more offers coming your way
20% OFF
Use code: COMEBACK
Tips & Trends
Create a space that's right for you
Seasonal Stories
New colors, patterns, and textures
We Haven't Seen You In A While
SHOP NOW
Fanatics
OFFICIAL PARTNER
COLLEGE
MLB
NFL
NBA
NHL
SALE
WE MISS YOU!
30% OFF
ORDERS OVER $50
Shop Now
Fanatics Rewards
on all purchases. Join For Free.
Look. We miss you. If you are unable to view this email, please click here.
Boden
WOMEN · MEN · KIDS · CLEARANCE
LOOK. WE MISS YOU.
...And here's what your wardrobe might be missing from us.

Not to sound too alarmist, but retail is evolving faster than ever before in the history of the known universe.

Brands that don't adapt—and marketers that don't adopt new technological skills—will not last.

One of those primary skills is systematically leveraging data to grow customer loyalty and improve customer retention.

This book focuses on retention, and how to build a churn prevention machine that keeps your customers engaged for the long run.

Introduction

Customers face an ocean of choices every time they make an online purchase.

30% of U.S. customers say they change brands often just for the sake of variety and novelty.[1] 49% of customers will gladly switch brands—for a coupon.[2]

In theory, when a brand acquires a new customer, that new customer is one flower in a thriving, growing garden of bountiful customer relationships. However, unlike flowers, customers can walk away at any time. Such a lack of loyalty poses a threat to every brand.

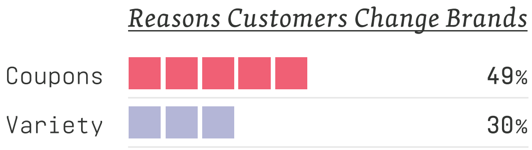

Acquiring customers is tough and expensive, but if those new customers don't bloom into loyal, repeat buyers, it can be extremely difficult for a retail brand to achieve sustainable growth.

Our last book explained how marketers are using customer analytics to significantly improve one-time buyer conversion rates (custora.com/oneandnotdone).

Now, we're focusing on reducing churn.

This book is an instruction manual for using the latest advances in customer analytics to build a scalable churn prevention machine.

Chapter 1

Churn Prevention Today

- What is customer churn?
- Why is it so important?
- What are brands doing about it?

What is churn?

The relationship you forge with a brand is just like *any* relationship. You meet someone. You hang out. If you enjoy each other's company, you'll hang out again. And again, and again, until a true relationship forms. Yet, over time, some friendships fade away. In some instances, we get caught up in our day-to-day lives and simply forget about each other. Sometimes a bad experience causes a rift. Sometimes, other relationships simply take priority.

Most of the factors that lead human beings to drift apart also apply to relationships between brands and customers. When a customer moves on from a brand, we call it *churn*.

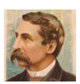

It's widely believed the word "churn" was first used in the context of customer attrition by Harrison Turnover McButterbean III in 1853. Also, churn prevention just sounds better than attrition prevention.

Nine times out of ten, once a relationship is over, it's over. And, if you've ever been dumped before, it's quite likely you spent a fair number of nights, weeks, and months thinking about "why?" and "what if?"

In hindsight, the signs were there for a while. She started sitting just a touch further away on the couch. He started to get off the phone a few minutes faster. She stopped surprising you with 'Swedish Fish' in your carry-on bag.

You wish you saw the signs then. Maybe there would have been a way to have a conversation. To change some behaviors. To get things on the right course before things faded away into a lukewarm, then lukecold state.

After reading this book, we want you, the marketer, to be able to detect the early warning signs of a customer break-up and to be prepared with an (automatic) intervention.

Most importantly, we want you to be able build a system to do that at scale.

Preventing just 1% of your very best customers from churning can lead to an overall revenue increase of 5%.

Why is churn prevention so important?

There's a reason brands are obsessed with customer churn: The economics of saving customers really adds up.

Repeat buyers are absolutely essential to long term growth. For a typical retailer, the top 10% of the customer base drives 50% of their revenue.

And every year, it's common to see half of those top customers lose some steam. If we can prevent 1% of that group from cooling down, that's the equivalent of adding 5% more revenue on the year. Yowza.

The Economics of Customer Churn[3]

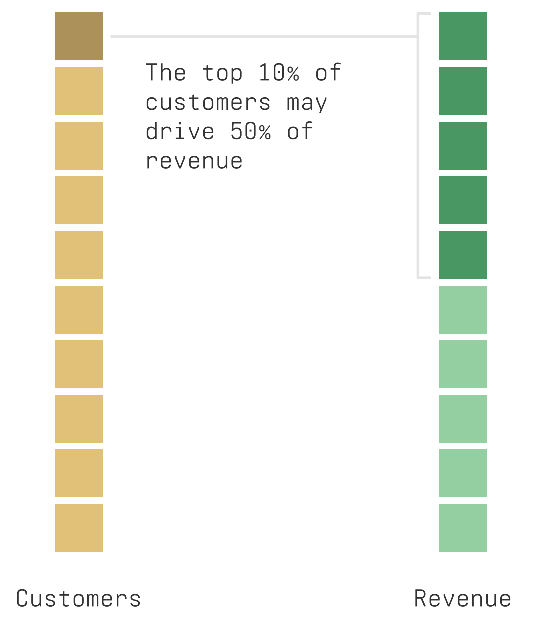

Churn is too important to ignore.

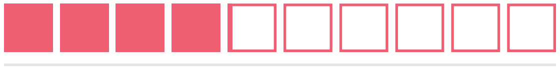

An average of **41%** of revenue comes from existing customers.[4]

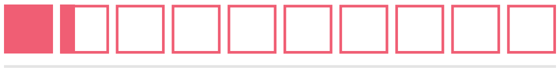

Only **13%** of customers are loyal to a brand.[5]

80% of brand “switchers” feel the company could have done something to retain them.[6]

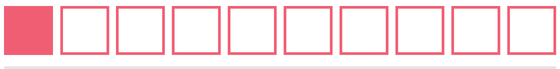

The top **10%** of customers often drive half of a retail brand's revenue.[7]

What are brands doing about churn?

In an ideal world, brands would be tuned in to the wants, needs, and feelings of every single one of their customers. In order to notice something as detailed as Swedish Fish disappearing from your carry-on, you have to be extremely familiar with the needs and desires of every customer.

This is a tall task. Think of your five closest friends. You know these people at an intimate level.

Jasper drinks hot coffee once at 10am, and again at 3:30pm, but he won't be caught dead with the hot stuff at 5pm or later.

Karen eats Chipotle every day. So much so that you wonder if you should say something. It's awkward.

These details make people who they are. And knowing these details—make relationships what they are.

Unfortunately, retailers have no practical way to tune into each relationship.

For one, customer data is not organized in a single place. For another, there simply isn't enough time to read through millions of individual customer stories.

As a result, most retailers rely on broad-stroke rules to guide their churn prevention program.

Rules aren't necessarily a bad thing. They can often be very helpful. No diving. No talking. No feeding after midnight. Stand on the right of the escalator, walk on the left. No seriously, stand on the right and walk on the left!

Rules can be great because they're simple and therefore easy to apply to large populations.

Most retailers have a rule-based program in place with primary rules such as "send a discount to anyone who hasn't made a purchase in 90 days."

From an operations perspective, getting a rule-based program off the ground is conceptually simple. Once a month, pull a list of customers who have crossed over that 90 day threshold. Send a "we miss you email" to that list.

These rule-based programs have two big flaws. One, the arbitrary, blanket "90 day" rule is often too late or too early for most customers. Two, though conceptually simple, even rule-based programs can be a real nuisance to manage.

Consider a fictional retailer, “Socktown,” that is using a one-size-fits-all, rule-based approach. Socktown's rule ensures that any customer who “has not made a purchase in the past 6 months” is considered churned.

Now imagine that Bob is a customer who shops every week— the guy simply loves socks. Some people wear socks down until they have holes in them. Not Bob.

He has a weird thing where he'll only wear a pair of socks twice before he moves onto a new pair. Bob is a special flower. It has something to do with putting socks in the wash, but none of us really understand.

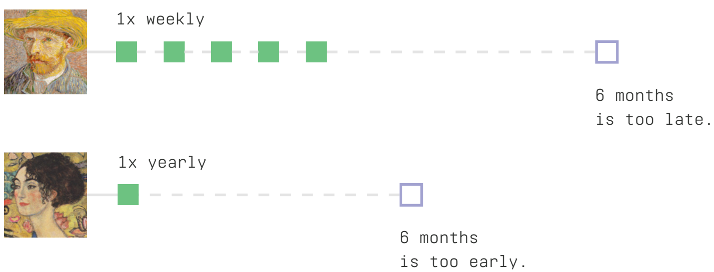

If Socktown waits until Bob is quiet for six months before reaching out with a churn prevention message, the opportunity to save the relationship will be long gone.

By that time, he may have moved on to buying all his socks at Amazon. We need a way to be much more responsive when someone like Bob shows signs of fading.

On the flip side, consider Anisha, who is a more typical sock shopper. She tends to load up with one big purchase in the holiday season. When she gets the 6-month email, she doesn't take action and make a purchase.

Instead she notes, “Ah, this is one of those retailers that will shower me with promotions if I don't order for a little while. Cool. I'll just always do that now, because you know, I like promotions.” Ideally, we'd only reach out to Anisha if we know she truly is fading off her typical habits — not sooner.

Rule-based programs sound simple, but when the rubber hits the road, marketers and analysts find they often require a surprising amount of time and elbow grease.

Does the scenario below sound familiar?

Hey, analytics team, love you all. Would you mind pulling a list of churning customers for me please?

Marketing Team

Sure, please submit a ticket and we'll add the request to the queue.

Analytics Team

------------------ Weeks go by ------------------

Here you go.

Analytics Team

They already churned! ¯_(^_^)_/¯.

Marketing Team

In the chapters that follow, we'll talk about a more productive, personal way to implement a churn prevention program.

Rather than relying on blanket 90-day rules, we'll use predictive analytics to detect when specific shoppers start veering off their norm.

We'll also discuss how to make programs easier to manage — so marketers can spend less time on logistics and more time thinking about how to connect with their fading-away customers.

Chapter 2

Blueprint for a Churn Prevention Machine

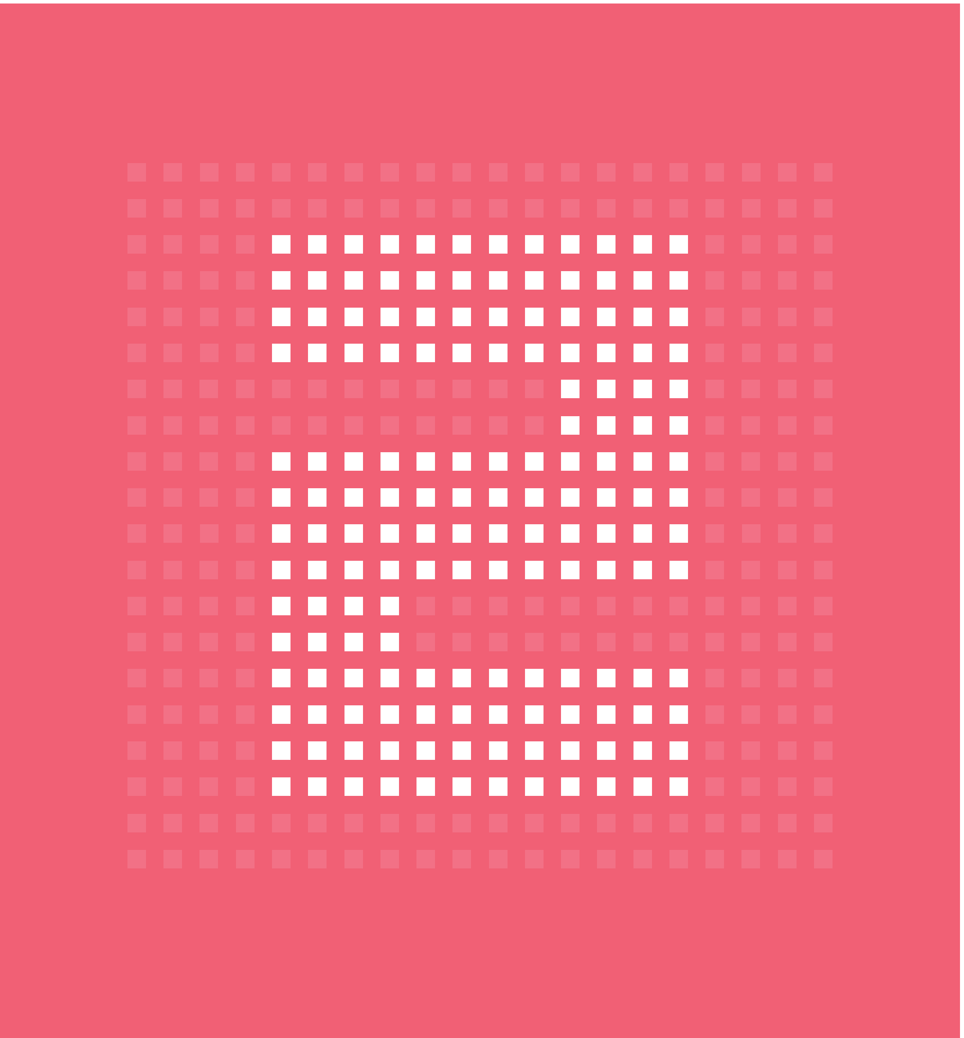

- Consolidate
- Predict
- Connect
- Test + Measure
- Automate

To move the needle on customer churn, we need to approach the problem with the mindset of an engineer. We need a system. A machine. Something that is always on, and something that works at scale. But this machine isn't simply an autopilot black box. Marketers need to drive it.

They need a way to monitor what is connecting with customers and what isn't. Ideally, the machine would handle the logistics of detecting who is at risk, connecting to campaign management tools, and measuring results, so that marketers can focus on the 'art' and save lapsing relationships.

A system like that has potential to make a difference.

So how do we get there?

While every organization is unique, we have found that most effective churn prevention systems share five core capabilities:

1. **Consolidate customer data**
2. **Predict behavior**
3. **Connect to marketing tools**
4. **Test and measure performance**
5. **Automate campaign delivery**

Whether you decide to build your system in-house or purchase an advanced customer segmentation platform (like Custora), you must have these five core capabilities if you want to maximize customer retention and minimize discounts.

This chapter will provide a deep dive into each of these five core capabilities.

The first two components concern data: getting (and keeping) it organized and extracting insights.

The base foundation of the machine we're putting in place should constantly be analyzing who is newly at risk for churn. That takes a lot of work.

The key to the first two components is being able to detect customer churn at the individual customer level.

Two things are needed for this. First, a holistic customer data profile that tracks the online and offline shopping behavior for every customer. Second, predictive analysis to create a churn "risk score" for each customer based on their behavior patterns.

Components three through five deal with applying those customers insights in the most effective ways to prevent churn.

We need an infrastructure in place that allows marketers to push campaigns to marketing channel tools in their technology stack, ("e.g.," an Email Service Provider or Facebook), run experiments, measure the incremental impact of the churn prevention program, and automate what's working.

Now let's take a closer look at each component.

1. Consolidate Customer Data

Customer data (including CRM data, transactional data, email data, on-site data, and more) is often siloed across data sources and marketing tools. With data spread across so many places, marketers struggle to build a complete view of their customers.

To address this challenge, you'll need to invest in a system or partner with a vendor that will connect the dots across your marketing technology stack.

MOBILE

IN-STORE

ONLINE

A word of caution: this is a large undertaking for an internal technology team to handle. With the scope and size of online and of offline data—and an endlessly growing pile of marketing tools with which to integrate — building requires substantial investment.

TRANSACTION

EMAIL

BROWSE

CUSTOMER RELATIONSHIP MGMT (CRM)

LOYALTY

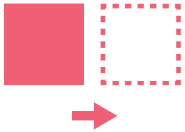

2. Predict

As we mentioned earlier, the key to saving a relationship is detecting early warning signs of churn.

This is where predictive analytics comes into play. Or, as the wind blows with popular terms of the season: Machine Learning, Artificial Intelligence, Fancy-Pants Statistics, Really-Good Math, Magic-Eight-Ball-Action, etc.

Whichever term you prefer, they all perform the same function: using algorithms that apply some form of pattern-matching technique to highlight customers who show signs of veering from their expected purchase behavior.

Vendors such as Custora can certainly lend a hand if you don't have the team (or if your team doesn't have the bandwidth) in house.

The best algorithmic approaches to detecting churn provide more than a simple binary “hot or cold” output. After all, that's not really how most relationships drift apart.

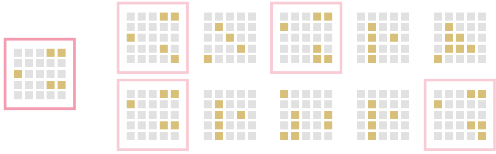

Rather than a simple “hot” or “cold,” some approaches will articulate where each customer stands in the stages of fading away.

In the case of our friend Bob (the guy who only wears socks twice and who usually buys a pair every week), a churn model might flag him as “at risk” if he misses two weeks. If he misses another week, we might think of him as “in serious trouble,” and so on.

Mathematically, statistical algorithms continuously assign a “probability of return” for Bob as more and more time passes since his last purchase. As that value decreases, he gets assigned into the appropriate churn stage.

Practically, this type of approach makes a lot of sense for the marketing team. We might think of different types of communication when Bob is just "cooling off" compared to when he is "highly at-risk."

Predictive modeling can offer more helpful insight as we think about how to reach out to Bob. For instance, models might surface that Bob is actually a high-value customer (one of those top 1% who drive so much value for the business!).

High-value customers often pay the biggest dividends for new churn prevention programs.

Knowing Bob's "VIP" status, we might be willing to roll out the red carpet and spend a bit more time and energy trying to prevent him from shifting his preference to Amazon.

Furthermore, advanced models might surface which brands and products each customer prefers. Perhaps Socktown happens to be losing all their "argyle" customers. The communication strategy for this group might vary substantially based on this information.

Predictive analytics can surface numerous other helpful insights, including discount preferences (e.g. free shipping or 10% off), channel preferences, and more.

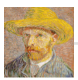

1+ PURCHASES BUT COOLING DOWN

Probability of Return: 50-60%

It's been three months since Bob has bought socks. It's a minor warning sign. Maybe he's working on his foot tan, or maybe he's moved on to Tevas? Or maybe he found a better sock shop. Whatever the case may be, he's defected from a consistent pattern. We call this "cooling down."

AT-RISK OF CHURN

(P)Return: 35-50%

Five weeks pass. Still no Bob. He's officially "at-risk" of churning.

HIGHLY AT-RISK

(P)Return: 20-35%

Ten weeks. No new socks. Something is definitely up. Bob is now "highly at-risk" of never making another purchase.

LOST

(P)Return: 0-20%

Fifteen weeks pass and it's safe to say that Bob is a "lost" customer. There is a very slim chance (less than 20%) that he's coming back.

3. Connect to Marketing Tools

Even if we have technology that surfaces the moment when specific customers start cooling down, we run into the old G.I. Joe mantra, "Knowing is only half the battle." We need a way for our marketing tools to ingest that info.

Most marketing teams utilize several different tools for customer communications. One tool to send email. Another for marketing on Facebook. Another for display ads. A different team or tool that manages direct mail. And so on.

Usually, if a marketing team creates a list of "at-risk" customers, they have to pass around a CSV file to all those teams and tools in order to get consistent messaging out the door. My Aunt Rhonda used to have a saying, "No one likes dealing with CSVs." She was right.

It's bad enough to have to deal with CSVs for a single segment of churned customers. The problem only gets worse when you start thinking about targeting segments within that list.

Approaching this with a systems mindset, we want things to work like clockwork. We want the segments of churning customers to be continuously connected with email, Facebook, and eventually all the other platforms in use across our team.

Churn prevention systems really take off when the segments of "at-risk" customers are automatically synchronized with each of those tools. Once they are, teams can set up automated programs to launch relevant, targeted campaigns on a weekly basis.

Now is a good time to pause and talk about “Farketing” and “Smarketing.”

So many marketing tools. So many customers. So much technology. So much to keep up with.

It's tough out there for a marketer. We know, because we work with some of the sharpest marketing teams on the planet.

And we spend a good portion of our time building things that will make their jobs easier and more efficient.

Overwhelmingly, the challenges we see that are facing marketers today boil down to two themes: Farketing and Smarketing.

Maybe you need more “Smarketing”?

Smarketing (Smarter Marketing) means that marketing leaders understand that their customer data has the potential to improve marketing results, but they struggle to unlock the value. It may be because they don’t have enough data scientists (or any?) in-house to keep up with the demand. Or possibly because they can’t afford to keep paying a CRM service provider to provide advanced segmentation. The net result is that the marketing team continues to miss opportunities to dramatically improve campaign performance.

Your primary challenge is generating insights.

Maybe you're challenged by “Farketing”?

Farketing* (Friction in the Marketing Department) means that the company has a top-notch data science team or agency, but the marketing team struggles to put those insights into action.

It could be due to lack of integration within the marketing technology stack. Or the time required to get segments built and files transferred.

But it means that marketing is missing opportunities due to friction in the process of using insights to drive better campaign performance.

Your primary challenge is scaling your insights to maximize the impact on the bottom line.

Maybe one of these challenges rings a bell? Maybe both? Either way, reading this book will help you on the path to more “smarketing” and less “farketing.”

**This term still hasn't caught on yet even though this is the second time we're including it in a book. We're still optimistic.*

4. Test and Measure Performance

We're now in a position where the data is organized, “at-risk” customers are continuously flagged, and all those customers will be shared across the various campaign management platforms in the marketing toolkit. The engine is revved up and it's now time for action.

But where to start? Though our marketer's intuition and customer analytics can help spark some communication ideas, we're unlikely to build a perfect outreach strategy on our first attempt.

The “art” side of the marketing equation takes time. Iteration. We need to embrace an experimental mindset.

We need to be willing to try a variety of ideas—from communication strategies to incentive packages—until we find which approaches connect with various customer types.

Experiment, measure. Experiment, measure. Experiment, measure.

You've probably heard the old carpenters' adage, “Measure twice, cut once.”

Around Custora, we have another saying. “The more you measure, the more you earn. Experiments turn the tide on churn.”

We actually have a number of similarly famous quotes.

“10% , 20% , 30% off, BOGO,
bye-bye margins, bye-bye CMO.”

“The email calendar is glorious and fun
the merch team saw it and now it's done.”

“Millennials don't care about what we say
Amazon eats all our margin away.”

But in all seriousness, measurement is an important topic. As we put our churn prevention engine in place, we need robust measurement tools to track our progress. We need a way to understand macro trends, as in, "are we moving the needle on customer churn?" And on the campaign-level, we need to know, "are the ideas we tried this week effective?"

Unfortunately, when it comes to measuring a customer metric such as churn, channel marketing teams offer little help. Most channel tools, such as ESPs, are built to measure immediate email-related metrics such as opens, clicks, or immediate conversions. If your marketing tactic motivates a customer to purchase something in a store, for example, your ESP will not be able to report on such a metric.

As marketers run campaigns across email, direct mail, Facebook, and display, it's important to track how many incremental customers are being saved.

This requires a system that makes control group testing easy.

Every week, as new segments of "at-risk" customers are detected, a small portion of customers should be isolated who won't be exposed to a promotion or offer.

Afterwards, you'll be able to compare revenue per user, profit per user, and overall conversion between the group who received your campaign and the group who didn't.

These metrics are crucial to track, but keep in mind that these represent only a fraction of the value being created and are only short-term metrics. The lifetime value of a customer who might have churned is significant.

Systems such as Custora can automate these sorts of metrics. In the absence of such a system, your internal CRM or Analytics team might also be able help out.

But, c'mon. Do we *really* need control groups?

Some marketers wouldn't even consider sending a churn prevention email campaign without a proper control group.

Others rarely use them in fear of "leaving money on the table."

The thinking goes, "I'm confident enough that this '15% Off' promo code we're sending to 10,000 customers at-risk of churn will generate additional sales (compared to doing nothing at all). Why would we risk lower revenues just to confirm what we already pretty much know will work?"

Ah. Well, surprisingly, the promo that we just *know* will outperform often doesn't. And it's important to know as soon as possible to ensure a timely correction.

Effectively managing control groups is especially important when churn prevention campaigns become automated. Often, that "free shipping" promo that consistently outperformed the control loses its luster after a few months.

5. Automate Campaign Delivery

Automation is essential for a successful churn prevention system.

As the team runs experiments and observes results on which messages are working with different types of customers, they need a way to set effective tactics on autopilot.

That's when the machine starts rolling. It's kind of like cruise control.

The machine is constantly reaching out to customers who show early signs of churn, those customers are getting the “Greatest Hits” of tactics from your experiments, and, as marketers, you occasionally turn off cruise control when ideas arise for things you'd like to try next.

Some of our customers talk about how, prior to having a machine like this in place, they would spend 20% of their time thinking of ways to connect with customers and 80% of their time dealing with logistics. After the machine is revved up, those numbers reverse. That's where the fun really kicks in. Not to mention the results!

Chapter 3

Operating The Churn Prevention Machine

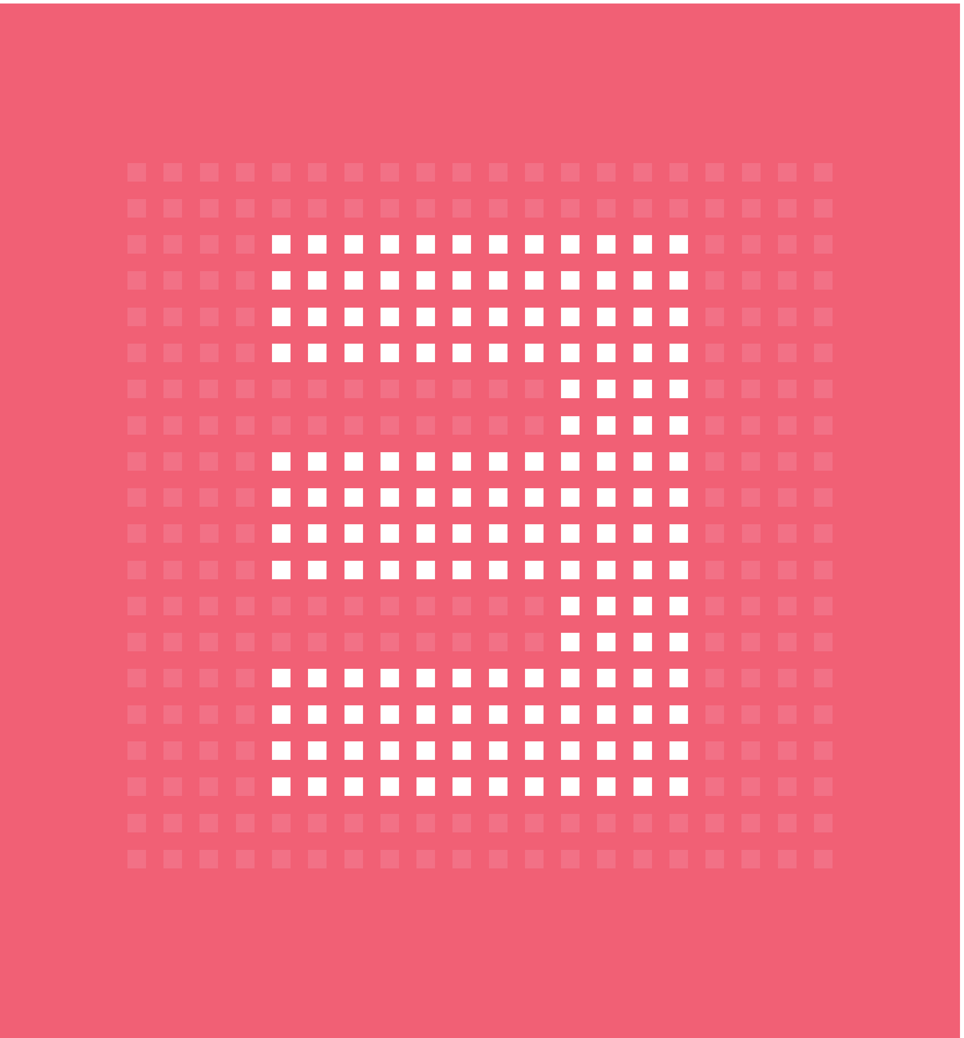

- Crawl
- Walk
- Run

Now that we've established the required components for a churn prevention system, it's time to put it to work.

You're not going to snap your fingers and suddenly have the world's best program in motion. The machine will enable you to get there, but as we mentioned, an effective churn prevention program requires a blend of art, science, and experimentation. It takes some time before things are fully running.

It's fun to imagine that a magic piece of technology will fix all your churn issues automatically as you set your feet up on your desk and secretly binge watch Netflix. Not quite. There are experiments to run, results to observe, and teams (creative and finance) with which to coordinate.

For this reason, Custora advocates a "crawl, walk, run" approach that starts simple and progressively becomes more complex as value is demonstrated.

We'll cover how to get quick wins ("crawl"), expand your program to drive more value ("walk"), and optimize efforts for maximum performance ("run").

Through each stage the program grows in two main ways:

1. **The number of sub-segments targeted within customers identified to be at risk of churn.**
2. **The number of marketing channels used.**

Crawl: Getting started

In the crawl stage, our goal is to identify customers in each predicted stage of churn and start winning some of those sweet sweet customers back.

To keep things simple, we suggest focusing on one channel in the "crawl" stage. Email is the most common starting point and often has an immediate impact on churn with the least amount of effort.

Now it's time to shift to the fun part: thinking of ways to connect with customers in those groups.

We won't roll out the fancy micro segments just yet. We want to test a handful of communication and promotional tactics for customers within each risk stage, see what works, then automate those effective ideas.

First, build an email experiment (with a control group).

The first thing we want to do is prepare for our experiment. We know we're going to test a handful of ideas as we reach out to customers across four stages of customer churn:

1+ PURCHASES BUT COOLING DOWN

AT-RISK OF CHURN

HIGHLY AT-RISK

LOST

The churn machine will take care of logistics like "who gets which message?" and "which ideas are working?" We just need to brainstorm which ideas to use for testing.

Keeping true to our crawl mantra, we believe it's most effective to start with the basics when it comes to messaging ideas as well.

When it comes to email, there's a variety of different messaging strategies to choose from. Depending on product/category mix, brand expectations, and regularly-scheduled email communications, churn prevention messages will vary in appropriateness and effectiveness.

Here are a few common categories of churn related messaging.

DISCOUNTS

% OFF

$ OFF

% OFF SPEND THRESHOLD

EARLY ACCESS TO PRODUCTS

SHIPPING UPGRADES

BUY-ONE, GET-ONE FREE (BOGOS)

Percentage discount-based programs afford simple yet effective experiment structures (and are quite successful for marketers setting up their first churn prevention machine), so let's start there.

The goal is to define a simple range of discount amounts. It's important to resist testing offers that are close together at this stage. A greater disparity will yield more informative results.

One popular tactic to help anchor the discount range is the following thought experiment:

If you knew that a high-value customer was on the verge of churning, what discount would you feel comfortable giving to save the relationship?

Your response represents the high end of the range and can help find the right starting point.

10% vs 25% off is a fine place to begin.

10% OFF (Low Discount)

25% OFF (High Discount)

NO DISCOUNT (Control)

And don't forget the control group.

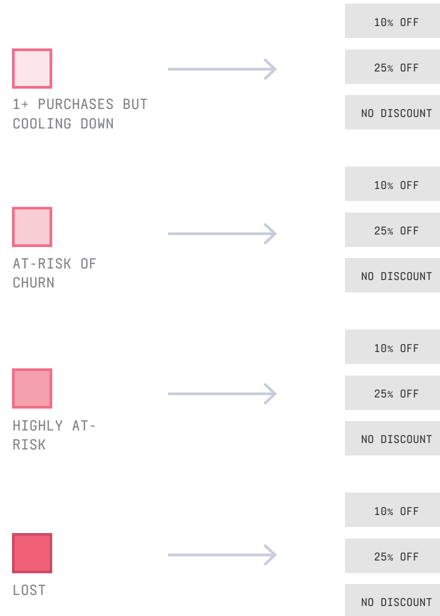

Okay, so we have three offers, and four segments of churn. Eeek—that's twelve different combinations to test.

This is where the value of the churn machine starts to sink in. We don't have to worry about setting up twelve experiments. We simply insert the three ideas into the mix, and the machine handles all the rest.

Ah but wait, we have the segments set up and we've got different offers to test, but what do we actually send?

We don't have time to create new artwork for dozens of different emails—we can barely keep up with our regular email creative as it is.

Again, it's okay to keep it simple to get started. Consider generic (but on-brand), text-based emails with a prominent offer if you decide to start with discounts.

With your experiment set up, it's time to hit "Send."

Because your churn machine is hooked up to your ESP, sending the campaigns is a snap. The machine will upload a list of customers who will receive one of the three promotional messages. Easy.

Next, after a few days, observe what worked.

It's time to review the campaign results. We need to see which customers actually came back and made a purchase.

Again, we're talking logistics now. Which means this falls into the territory of things our machine makes easy. Rather than waiting for our internal analytics team to produce some results, the machine will dish out reports that look like the following:

Churn Prevention Email Experiment Profit Per User

Lifecycle Stage	No Discount Message	Low Discount (10%)	High Discount (25%)
1+ Purchases but Cooling Down	$0.05	$0.36	$0.27
At-Risk	$0.03	$0.22	$0.19
Highly At-Risk	$0.01	$0.11	$0.14
Lost	$0.01	$0.07	$0.12

Winning message

Since we're talking about churn, it's important to examine profit per user. It's tempting to get caught up in Return on Ad Spend (ROAS), last-click revenue and overall revenue.

But we're testing discount strategies here, so it's imperative that we measure on a profit basis. After all, if you really want to "save" customers, you could hand out 99% off discounts and you'll see those retention rates go through the roof!

In the “one or more purchases but cooling down” recipient group, the 10% off “low discount” was most effective.

For the “at-risk” group, the low discount also won.

However, in the “highly at-risk” and “lost” groups, the high discount drove the highest incremental revenue.

Unsurprisingly, different ideas have varying effectiveness across different churn stages. It required a larger “lure” to attract back the least engaged customers.

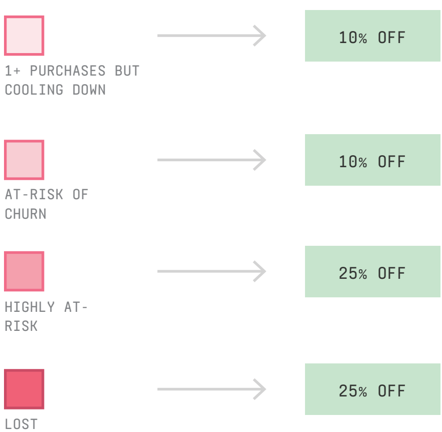

For the “cooling down” customers, the low discount appears to be most effective, though it might be the case that the high and the low discount each drove the same conversion rate.

In that case, we generated more profit per user with a lower discount, so that's the winning message we'll continue to send to customers in this stage.

During latter stages of churn, the small discount simply didn't move the needle and the larger one was required.

Finally, automate the most effective messages.

Now that we know which offers worked within each churn segment, the machine can automatically detect who's in which stage each week, export those lists to the ESP, and send the optimal offer on an ongoing basis.

We'll also keep an eye on the results and control groups in case any offer or messaging adjustments are needed.

Walk: Adding customer value into the mix

With a basic system in place and running smoothly, it's time to kick it up a notch.

The goal of the “walk” stage is to get even more granular with our segmentation. Now that we're already segmenting by churn stage, the next step is further segmenting by value tier.

We'll focus on the churning customers with the highest probability of return (the “cooling off” shoppers) and sub-segment that group by three tiers of value: high-spending VIPs, medium value shoppers, and lower-value, discount-oriented shoppers.

We'll also expand into an additional marketing channel (Facebook) to reach and win back even more customers.

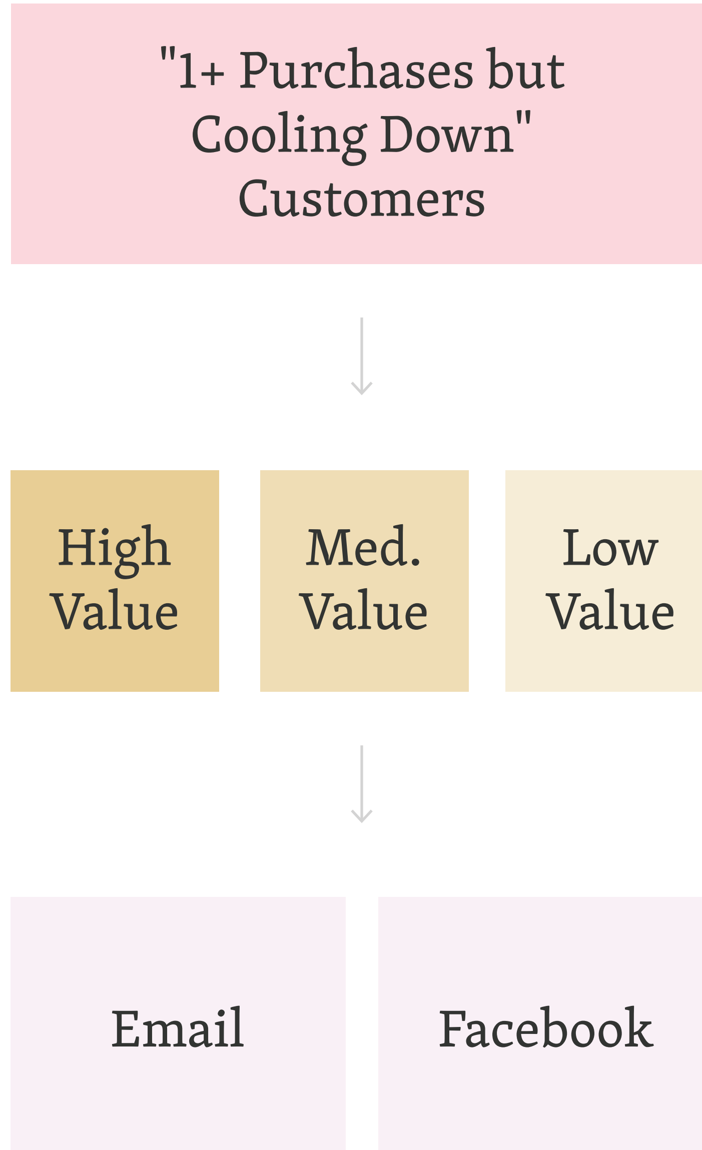

The basic churn system ensures that we're reaching out to customers at the right time and with the right promotional tactic. However, there's a lot more we can do.

Remember our two sock shoppers, Bob and Anisha? They both love socks, but have very different shopping habits and preferences.

Bob is hugely valuable to the brand, whereas Anisha is more a typical, seasonal shopper.

While Socktown wants to retain both customers, it's critical for Socktown to roll out the red carpet to save their relationship with a high-value customer like Bob.

Segmentation by Customer Value Tier

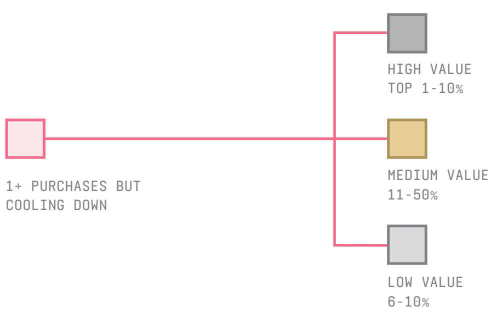

Remember, the top 10% of a customer base can contribute more than 50% of the revenue.

In the walk stage, we want to start making the communication more relevant within each churn stage.

Our experience has shown that the first way to segment further is to break down by customer value tier withing the cooling off stage. Instead of a one-size-fits-all approach, we can think about treating those VIP customers differently.

In addition to adding value tier segmentation, we can also start experimenting with targeted offers on Facebook. Facebook impressions are expensive, but if they save those high-value customers, it's worth it!

Let's have another look at our campaign results, but broken down by customer value.

Churn Prevention Email Experiment Results

Profit per user for "Cooling Down" Customers broken down by Customer Value Tier

Lifecycle Stage	No Discount (Control)	Low Discount (10%)	High Discount (25%)
1+ PURCHASES BUT COOLING DOWN	$0.05	$0.36	$0.27
HIGH VALUE TOP 1-10%	$0.04	$0.02	$0.03
MEDIUM VALUE 11-50%	$0.05	$0.43	$0.25
LOW VALUE 6-10%	$0.03	$0.25	$0.22

Winning message

Iiiinnteresting. Although the 10% off message was most effective overall for the "cooling-down" segment, when we slice by value, we discover the highest profit per VIP customer was from the "no discount" control message. And that $.04 per user was far below any other winning result. What does this mean?

Basically, our discount-based email experiment had no effect on the top value-tier customers showing the first signs of churn. Bummer.

So what can we do?

First, let's build a new experiment to test an alternate value proposition or call to action for these special customers. Maybe it's a short and sweet letter from the CEO expressing how much they are appreciated and encouraging them to check out they latest releases they may have missed since they've been gone.

The new test might look something like this.

New Churn Prevention Experiment Targeting High Value Customers

Lifecycle Stage	No Discount (Control)	Low Discount (10%)	Letter from CEO
1+ PURCHASES BUT COOLING DOWN			
HIGH VALUE TOP 1-10%			

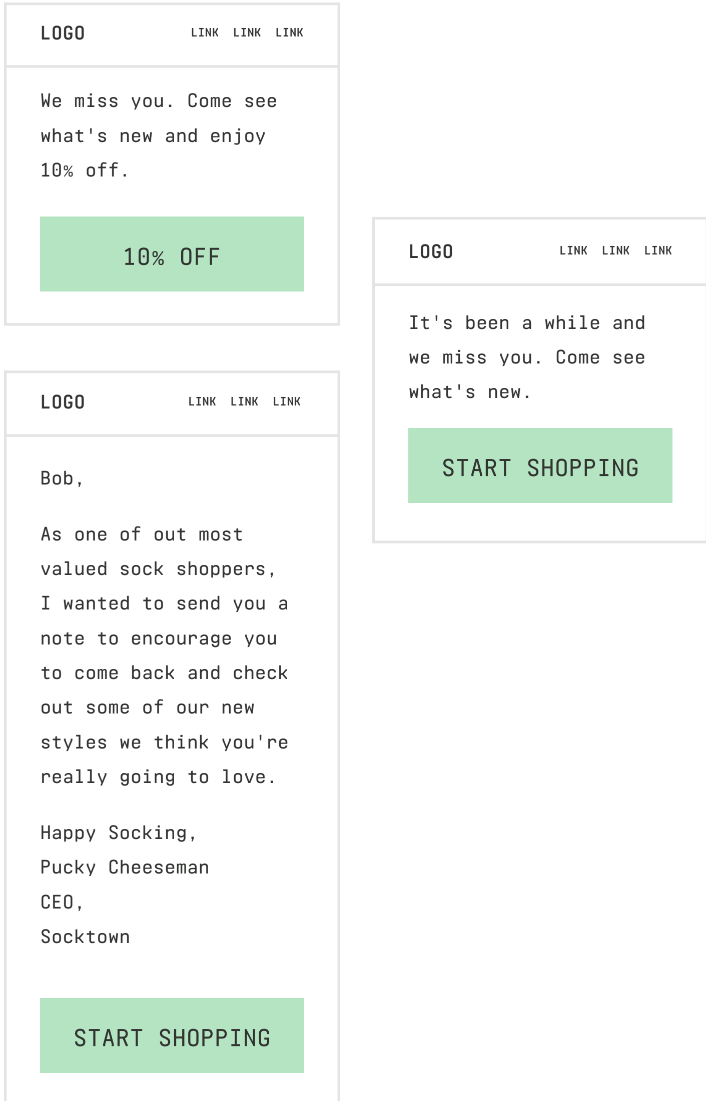

We'll keep an eye on that campaign and hopefully see an improvement in profit per user. Maybe a letter from Pucky is the nudge Bob needs to get back to his socking ways.

Next, let's expand beyond email into an additional marketing channel. Facebook is a great place to start —especially when your churn machine is set up to export weekly customer lists directly into Facebook's Custom Audiences ad platform.

Before uploading the "cooling down" VIP list to Facebook, we'll want to figure out our communication strategy. Again, we suggest starting simple. Focus on a few, popular products and a clear call to action. Maybe mention new best sellers that might reignite the flame for the waning shopper.

Remember, the big improvements here are:

1. Reconnecting with customers that you know are slipping away.
2. Finding the right message and channel combination.

Here's what our entire set-up for the "walk" stage might look like.

"Walk": Message & Channel Distribution

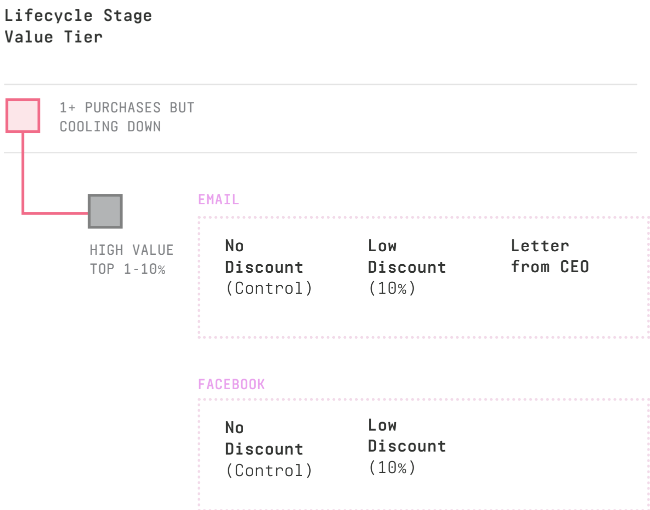

With our "cooling-down" VIPs now receiving communications across a few channels, we've increased the likelihood of winning them back.

And with everything set up on a weekly recurring basis, we're building more momentum towards having a significant, sustainable impact on customer churn.

Next, let's pick up the pace with a new sub-segment and marketing channel.

Run, Marketer, Run.

So far, we've got weekly segmentation by predicted churn stage and have started breaking down those "one purchase but cooling down" customers by customer value. We've also expanded communications across email and Facebook.

Within a few weeks, there should be an evident impact on churn. Excellent.

It's time to layer in some more nuanced behavioral segmentation, "personas," and add a new channel, direct mail.

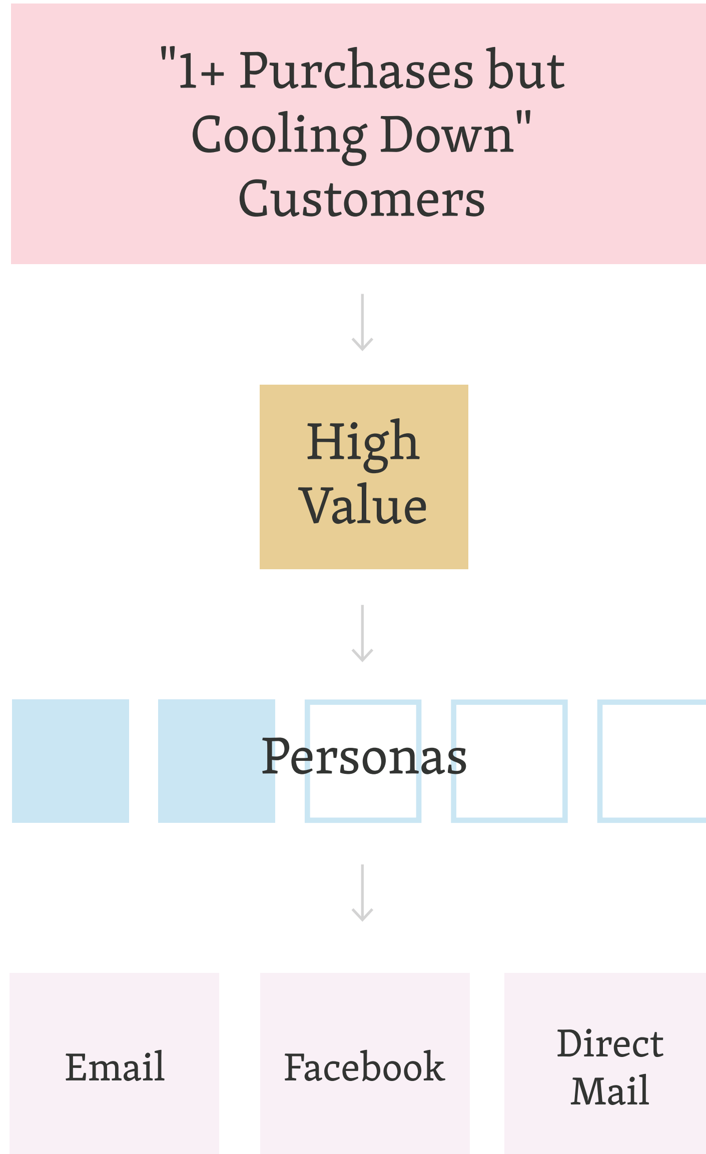

These aren't your ad agency's personas.

Remember, your churn machine employs advanced statistical capabilities, so let's see what else we can discover to further inform our churn prevention communications.

One technique known as "cluster analysis" segments customer populations into groups based into distinct purchasing patterns and demographics.

Whereas traditional marketing "personas" might rely on intuition and depict qualitative generic (but disparate) demographic portraits, "cluster analysis" asks, "what does the data reveal?"

Additionally, marketing personas tend to be historical in nature (e.g., what have we seen from our customers?) as compared to predictive personas (e.g., what might we expect?).

Using our Socktown retailer as an example, let's further sub-segment our high-value, "cooling-down" customers and see what personas we find.

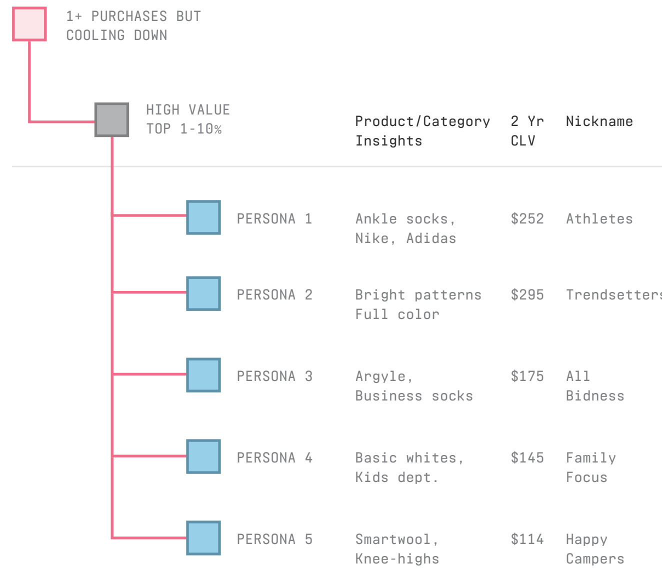

Personas 1 ("Athletes") and 2 ("Trendsetters") have the highest CLV and represent 23% and 20% of the cooling down VIP segment, respectively. So again, keeping with the spirit of simplicity, let's start running with the two most important personas. We can always come back and expand our communications.

Time to get a little more creative.

Thus far, we've leveraged fairly generic creative assets and templates.

With our new behavioral insight into different personas, we have the opportunity to send more nuanced and personalized communications.

This means you'll need some help from the creative team.

Ah, creative. The elephant. The bottleneck. As the saying goes, “Creative is where campaign ideas go to die.”

Okay, that's a little harsh. No one probably ever said that, but there's a good chance it's crossed a marketer's mind.

The reality is, marketers need creatives. The creative team plays a critical role in crafting communications that will truly connect with customers.

Unfortunately, producing numerous, high-quality creative also requires manual time and effort.

Generally speaking, “personalization” at scale and “custom creative” don't go hand-in-hand.

Fortunately, when it comes to setting up a churn prevention system, your creative can be “evergreen.” What do we mean by “evergreen”?

You can create a small number of creative assets and reuse them again and again.

Because each customer passes progressively through each stage of churn, we don't have to worry about the same customer seeing the same message twice.

This means we can create one robust set of creatives and use it for quite a while. And using a “Crawl, Walk, Run” approach gives us the chance to learn what’s working and build as we discover more.

Getting back to our churn prevention campaign focused on the most important personas within the "cooling down" VIP segment, here's how our latest experiment might shape up.

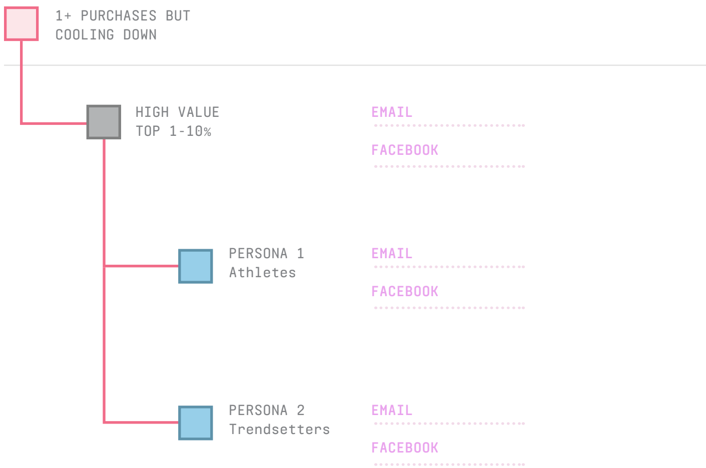

After setting up a few new discount-based experiments with custom creatives for our personas in both Facebook and email, we can also move into direct mail.

Direct mail, once a dominant marketing channel for many retailers, can still be very effective when approached with the analytical mindset of a churn prevention engineer.

Yes, it's crazy expensive sending a bunch of dead trees to thousands of shoppers, but when you can identify valuable customers and their unique preferences, the return per send can shift dramatically.

Tiffany & Co., for example, saw a 17% lift in revenue per catalog recipient after re-focusing their distribution to high-value customers. Yowza.

New vendors, such as PebblePost, are making it easier than ever to connect with customer data platforms and fully leverage the potential of more personalized mailers.

As always, it's important to monitor results and make adjustments as the effectiveness changes across campaigns and channels.

We're running, we're sweating, and we're seeing a big impact on churn. And we've only just covered a fraction of churn stages, value tiers, and customer personas. Here's a look at how far we've come.

Blueprint for a Customer Churn Prevention System

Moving from basic to advanced with predictive customer segmentation and multiple marketing channels

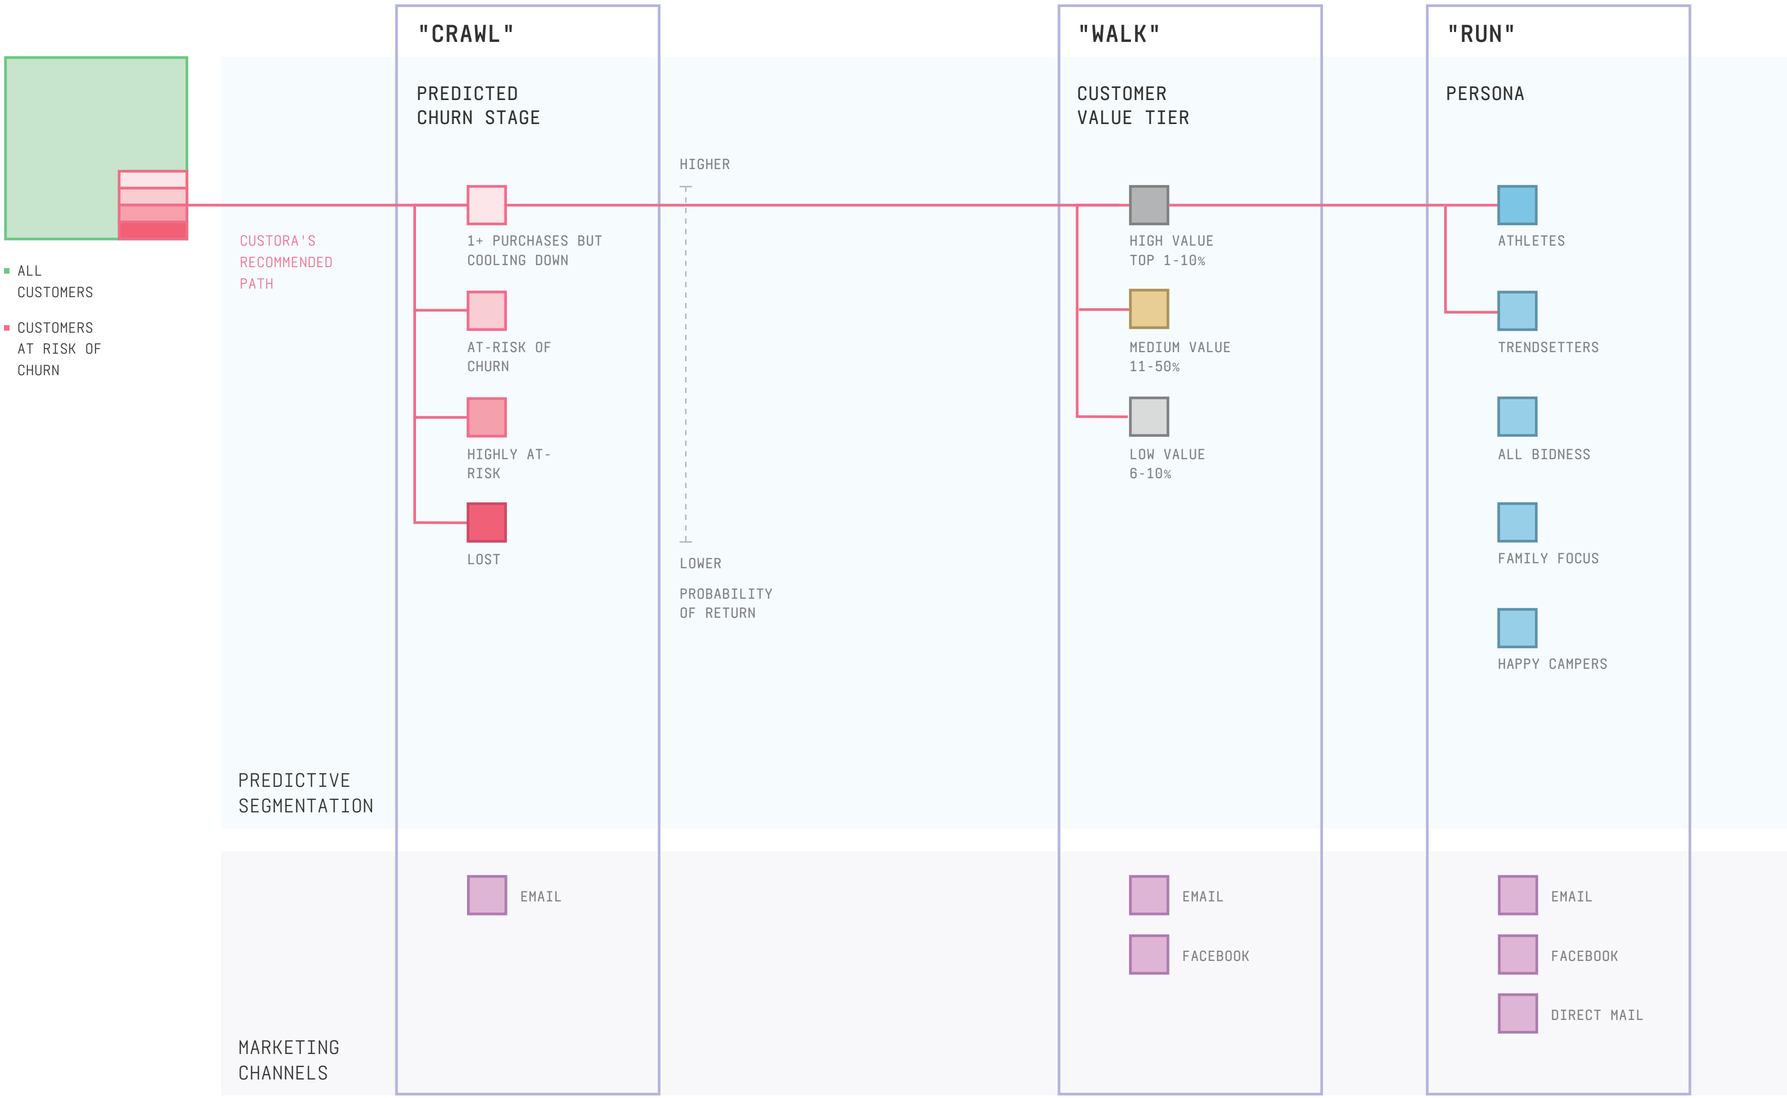

Looking ahead, a retailer with five robust personas might want to develop five versions of their email at each stage in the churn lifecycle.

This introduces additional overhead to build and maintain, but can often lead to meaningful improvements in churn prevention performance.

Segmenting on both personas (for creative optimization) and value segments (for offer/call to action optimization) is the sign of an advanced churn prevention system.

We should continue to refresh email creative periodically (e.g., quarterly) and look for additional optimization opportunities, but at this point, it's time to transition the program into an always-on, steady-state mode and look to the next set of incremental opportunities.

Chapter 4

Conclusion

Retailers invest tremendous marketing resources into acquiring new customers. But it's a losing game when those customers don't stick around and make repeat purchases.

The key to fundamentally transforming customer economics is to cultivate long-term relationships with customers from the moment they first engage with the brand and to step in with the right message when they show signs of fading away. This requires an always-on, steady-state churn prevention system.

Custora can help. The platform provides the data aggregation and predictive analytics to monitor which customers are at risk of churn, and integrates with tools utilized across the marketing team to make it easy for everyone to take action on those insights.

If you're getting serious about reducing customer churn, we'd love to chat.

About Custora

Custora is an advanced customer segmentation platform.

We help marketers improve the ROI of their email, display, direct mail, and Facebook campaigns. Our platform gives marketers the power to create smart customer segments and activate them everywhere.

Custora aggregates customer data from disparate sources, leverages predictive analytics to surface optimal segments, integrates with marketing channels, campaign tools, and personalization technologies for easy segment activation, and measures the impact of each campaign at the customer level.

Contributors

Matt Achariam

Carrie Battan

Maddie Buras

Jordan Elkind

Nicole Fleming

Martin Gordon

Wendy Hu

Andrew Lim

Corey Pierson

Taylor Simpson

Sara Skov

David Stychno

Kevin Young

References

[1, 2] Retaildive.com

[3] Custora

[4] Adobe Loyal Shoppers Report 2013

[5] McKinsey & Company

[6] Accenture

[7] Custora